

CS 483 Project Report

Team_name: hawks

Names: Nikhil Paidipally, Daniel Niewierowski, Nicolas Nytko

Net_IDs: npaidi2, nnytko2, dniewi2

RAI_IDs: 5d97b1f588a5ec28f9cb94b0, 5d97b1f488a5ec28f9cb94ae,
5d97b1f388a5ec28f9cb94ac

Affiliation: uiuc

MILESTONE 1:

Kernels that collectively consume more than 90% of the program time:

- [CUDA memcpy HtoD]
- volta_scudnn_128x64_relu_interior_nn_v1
- volta_gcgemm_64x32_nt
- fft2d_c2r_32x32
- volta_sgemm_128x128_tn
- op_generic_tensor_kernel
- fft2d_r2c_32x32

CUDA calls that collectively consume more than 90% of the program time:

- cudaStreamCreateWithFlags
- cudaMemGetInfo
- cudaFree

Difference between kernels and API calls:

CUDA kernels are effectively like regular C functions that are executed an arbitrary number of times in parallel by an arbitrary number of different CUDA threads on the gpu/device.

CUDA API are effectively regular C functions that execute on the cpu/host and they are not executed in parallel like the CUDA kernels. Some examples of CUPA APIs include cudamalloc, cudaMemCpy, cudaFree, etc.

Rai running MXNet on the CPU:

(CPU C Code)

Loading fashion-mnist data... done

Loading model... done

New Inference

Op Time: 11.330177
Op Time: 74.542234
Correctness: 0.7653 Model: ece408
100.71user 11.16system 1:29.81elapsed 124%CPU (0avgtext+0avgdata
6044220maxresident)k
0inpu
ts+0outputs (0major+2308767minor)pagefaults 0swaps

(CPU Python Code)

Loading fashion-mnist data... done
Loading model... done
New Inference
EvalMetric: {'accuracy': 0.8154}

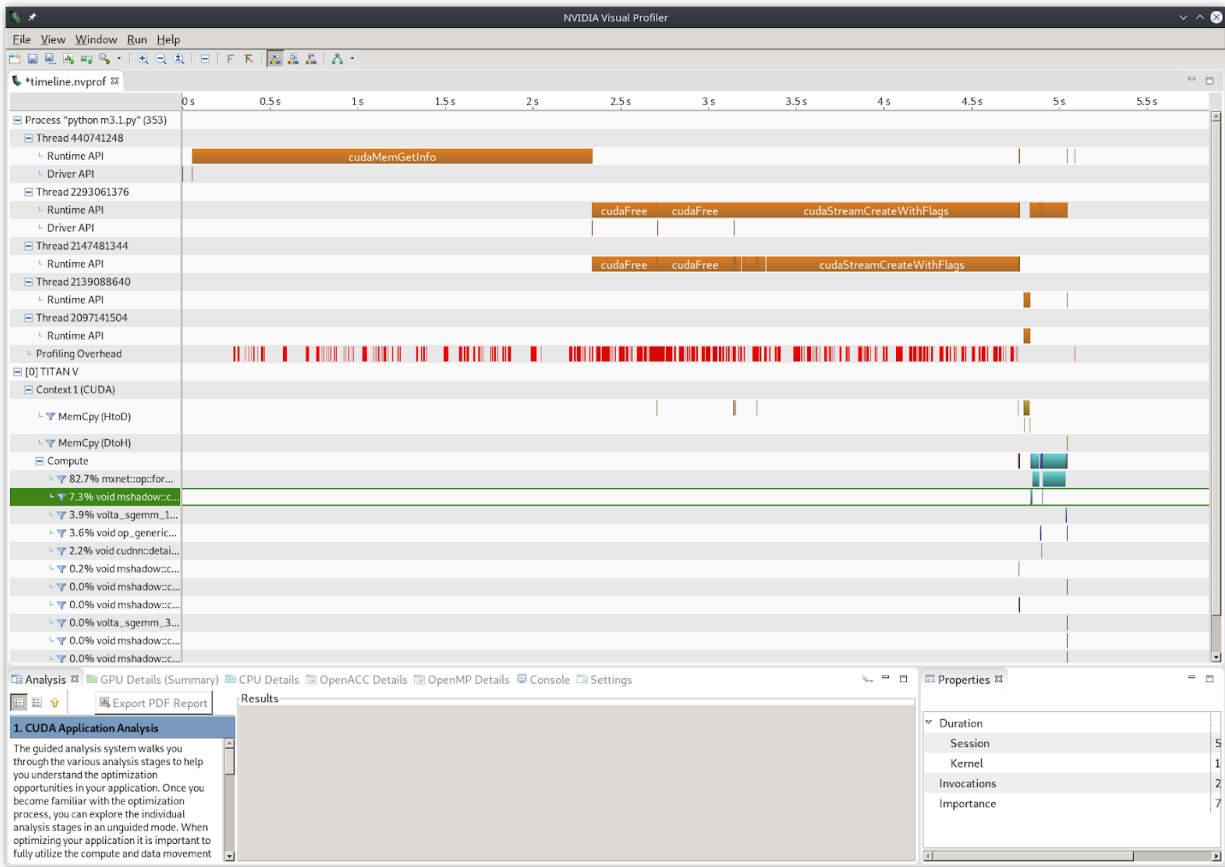
17.90user 4.36system 0:09.78elapsed 227%CPU (0avgtext+0avgdata
6045072maxresident)k
0inputs+2824outputs (
0major+1603748minor)pagefaults 0swaps

Rai running MXNet on the GPU:

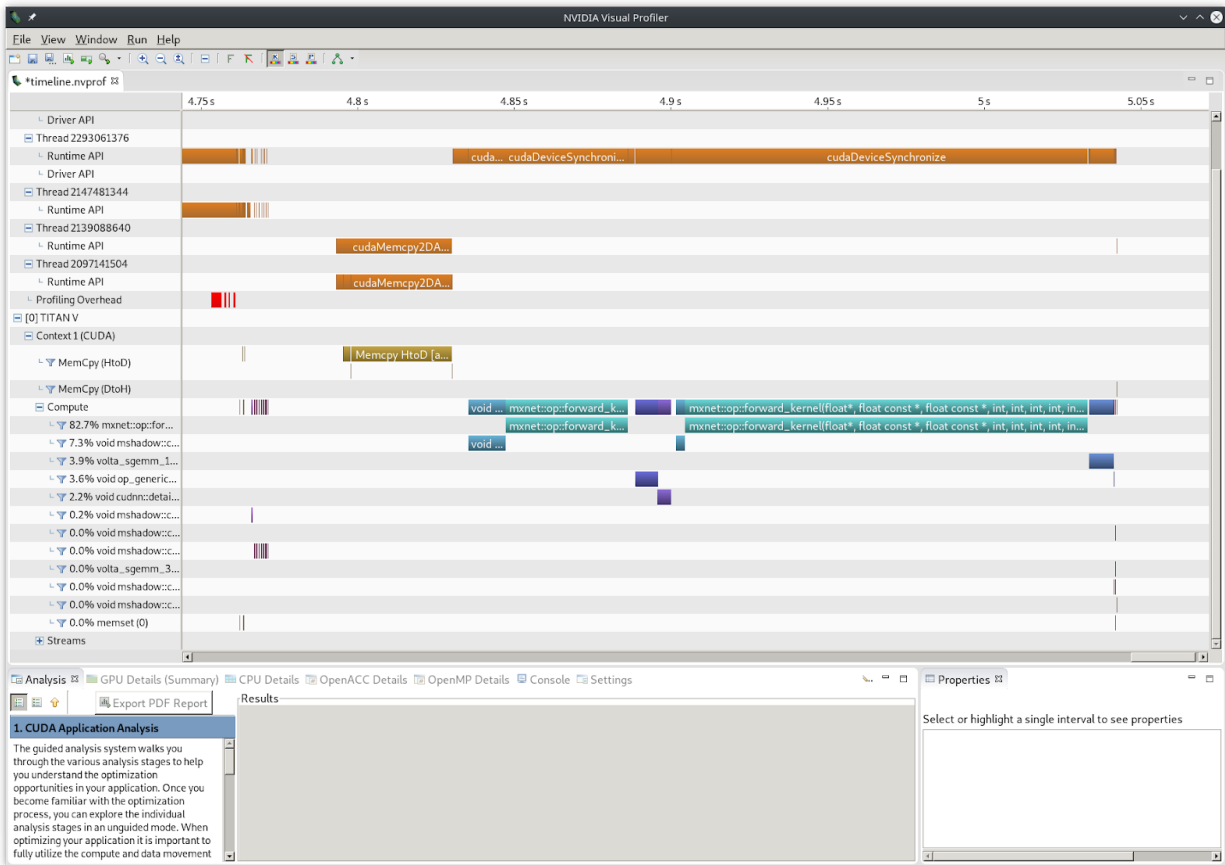
Loading fashion-mnist data... done
Loading model... done
New Inference
EvalMetric: {'accuracy': 0.8154}

4.96user 3.07system 0:04.63elapsed 173%CPU (0avgtext+0av
gdata 2996488maxresident)k
0inputs+1712outputs (0major+734016minor)pagefaults 0swaps

MILESTONE 3:



(View of overall CUDA execution)



(Zoomed view of the forward-prop convolution kernel)

MILESTONE 4:

To establish a baseline, we first ran our convolution without any optimisations several times and recorded the average running time for both operations, detailed below:

Op Time: 0.045837

Op Time: 0.134574

Correctness: 0.755 Model: ece408

This was our “time to beat” so to speak.

OPTIMISATION 1:

The first optimisation we performed was to tune the kernel with the restrict keyword unroll the innermost loop. The `__restrict__` keyword was added to the arguments y, x, and k to signal to the compiler that the memory regions are non-overlapping and thus the pointers do not alias. In some cases this can provide significant speed-ups. The innermost convolution loop was also unrolled, since the mask size is fixed at 5 elements in each direction.

Implementing the above improved our times slightly, but it was hard to judge overall effectiveness due to the randomness of the timing values. Perhaps the loop control does not affect the running time too heavily.

Op Time: 0.036734
Op Time: 0.111906
Correctness: 0.755 Model: ece408

OPTIMISATION 2:

Our second optimisation was to improve the kernel itself by using shared memory to cache tiled sections of the image. For each image channel, a block of threads will first load in a tiled section together, then perform the convolution on the tiled section. This approach was modeled after the “strategy 2” tiled convolution proposed in lecture.

Additionally, we were able to catch a bound error and increase our correctness slightly.

Op Time: 0.026975
Op Time: 0.076901
Correctness: 0.7643 Model: ece408

OPTIMISATION 3:

For our third optimisation, we performed a sweep of the block values in order to determine the best (local) running time. Since the maximum block size is defined as 1024 threads in CUDA, our domain of block values can range from 5 to 32.

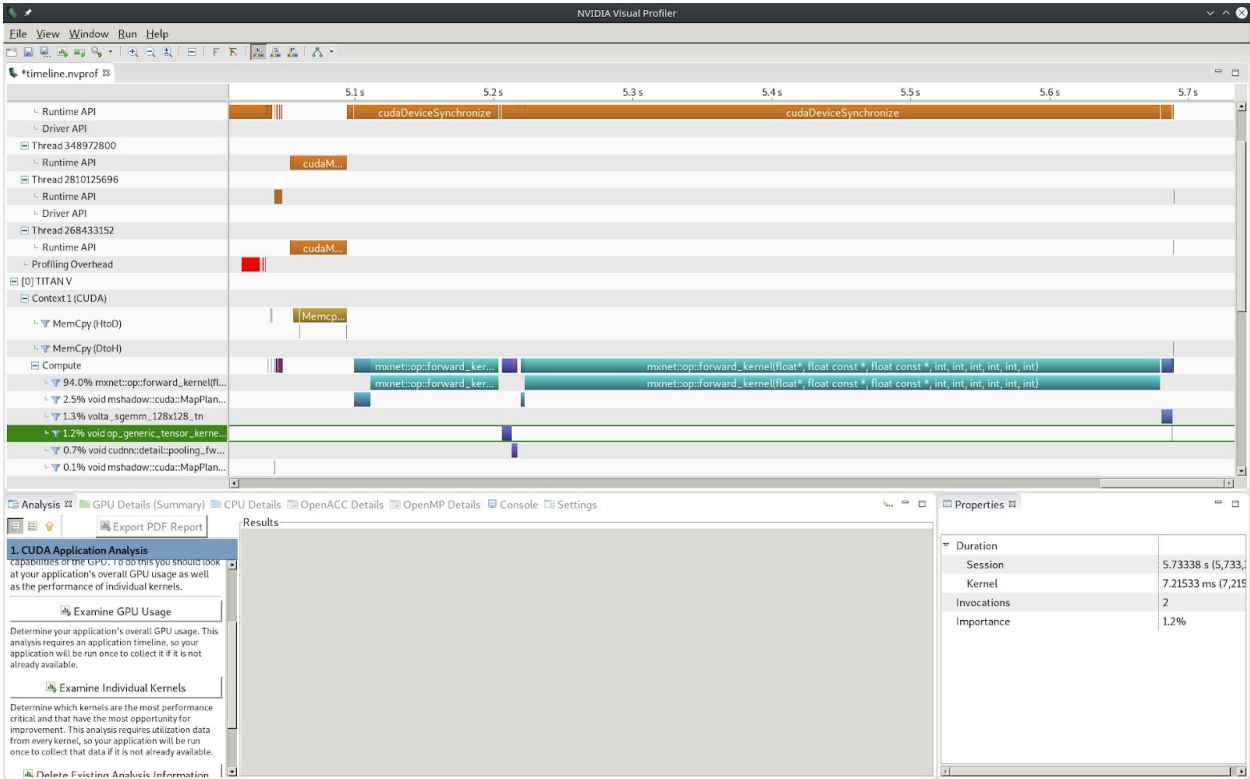
We modeled our sweep after the golden section search, a method for optimising 1-dimensional functions.

Block Size	Op1 Time (s)	Op2 Time (s)
15	0.09702	0.533354
21	0.08562	0.456664
25	0.123988	0.681743
19	0.109751	0.39073
20	0.131026	0.437012

22	0.9062	0.507769
----	--------	----------

The local minimum was found experimentally to be a block size of 21.

PROFILER VIEW:



(View of the optimised kernel in NVVP)